

PRODUCTION READY MICROSERVICES AT SCALE

RAJEEV N B
GO-JEK

RAJEEV N BHARSHETTY

- ▶ Product Engineer at GO-JEK
- ▶ Distributed Systems | Security | Data
- ▶ Building reliable and scalable systems
- ▶ **@rbharshetty** - Twitter
- ▶ **@rshetty** - Github

TRANSPORT, LOGISTICS, HYPERLOCAL DELIVERY AND PAYMENTS

- ▶ 18+ products
- ▶ 1 Million+ Drivers
- ▶ 500+ Microservices
- ▶ 15k+ Cores
- ▶ 2 Cloud Providers
- ▶ 6 Data centres
- ▶ 100 Million+ bookings a month



GO-JEK expands to 4 new markets - Thailand, Vietnam, Singapore & Philippines

AGENDA

- ▶ What are **Production Ready** Microservices ?
- ▶ Why do we need them ?
- ▶ How do we build them ?
- ▶ Future work
- ▶ Conclusion

MICROSERVICES

**SMALL, AUTONOMOUS
SERVICES THAT WORK
TOGETHER**

Sam Newman
(Building Microservices)

**PRODUCTION
READY**



BUILDING TRUST

WHY ?

Goal is to be **Available** to serve our
users

HOW ?

PRODUCTION READINESS CHECKLIST

WHY CHECKLIST ?

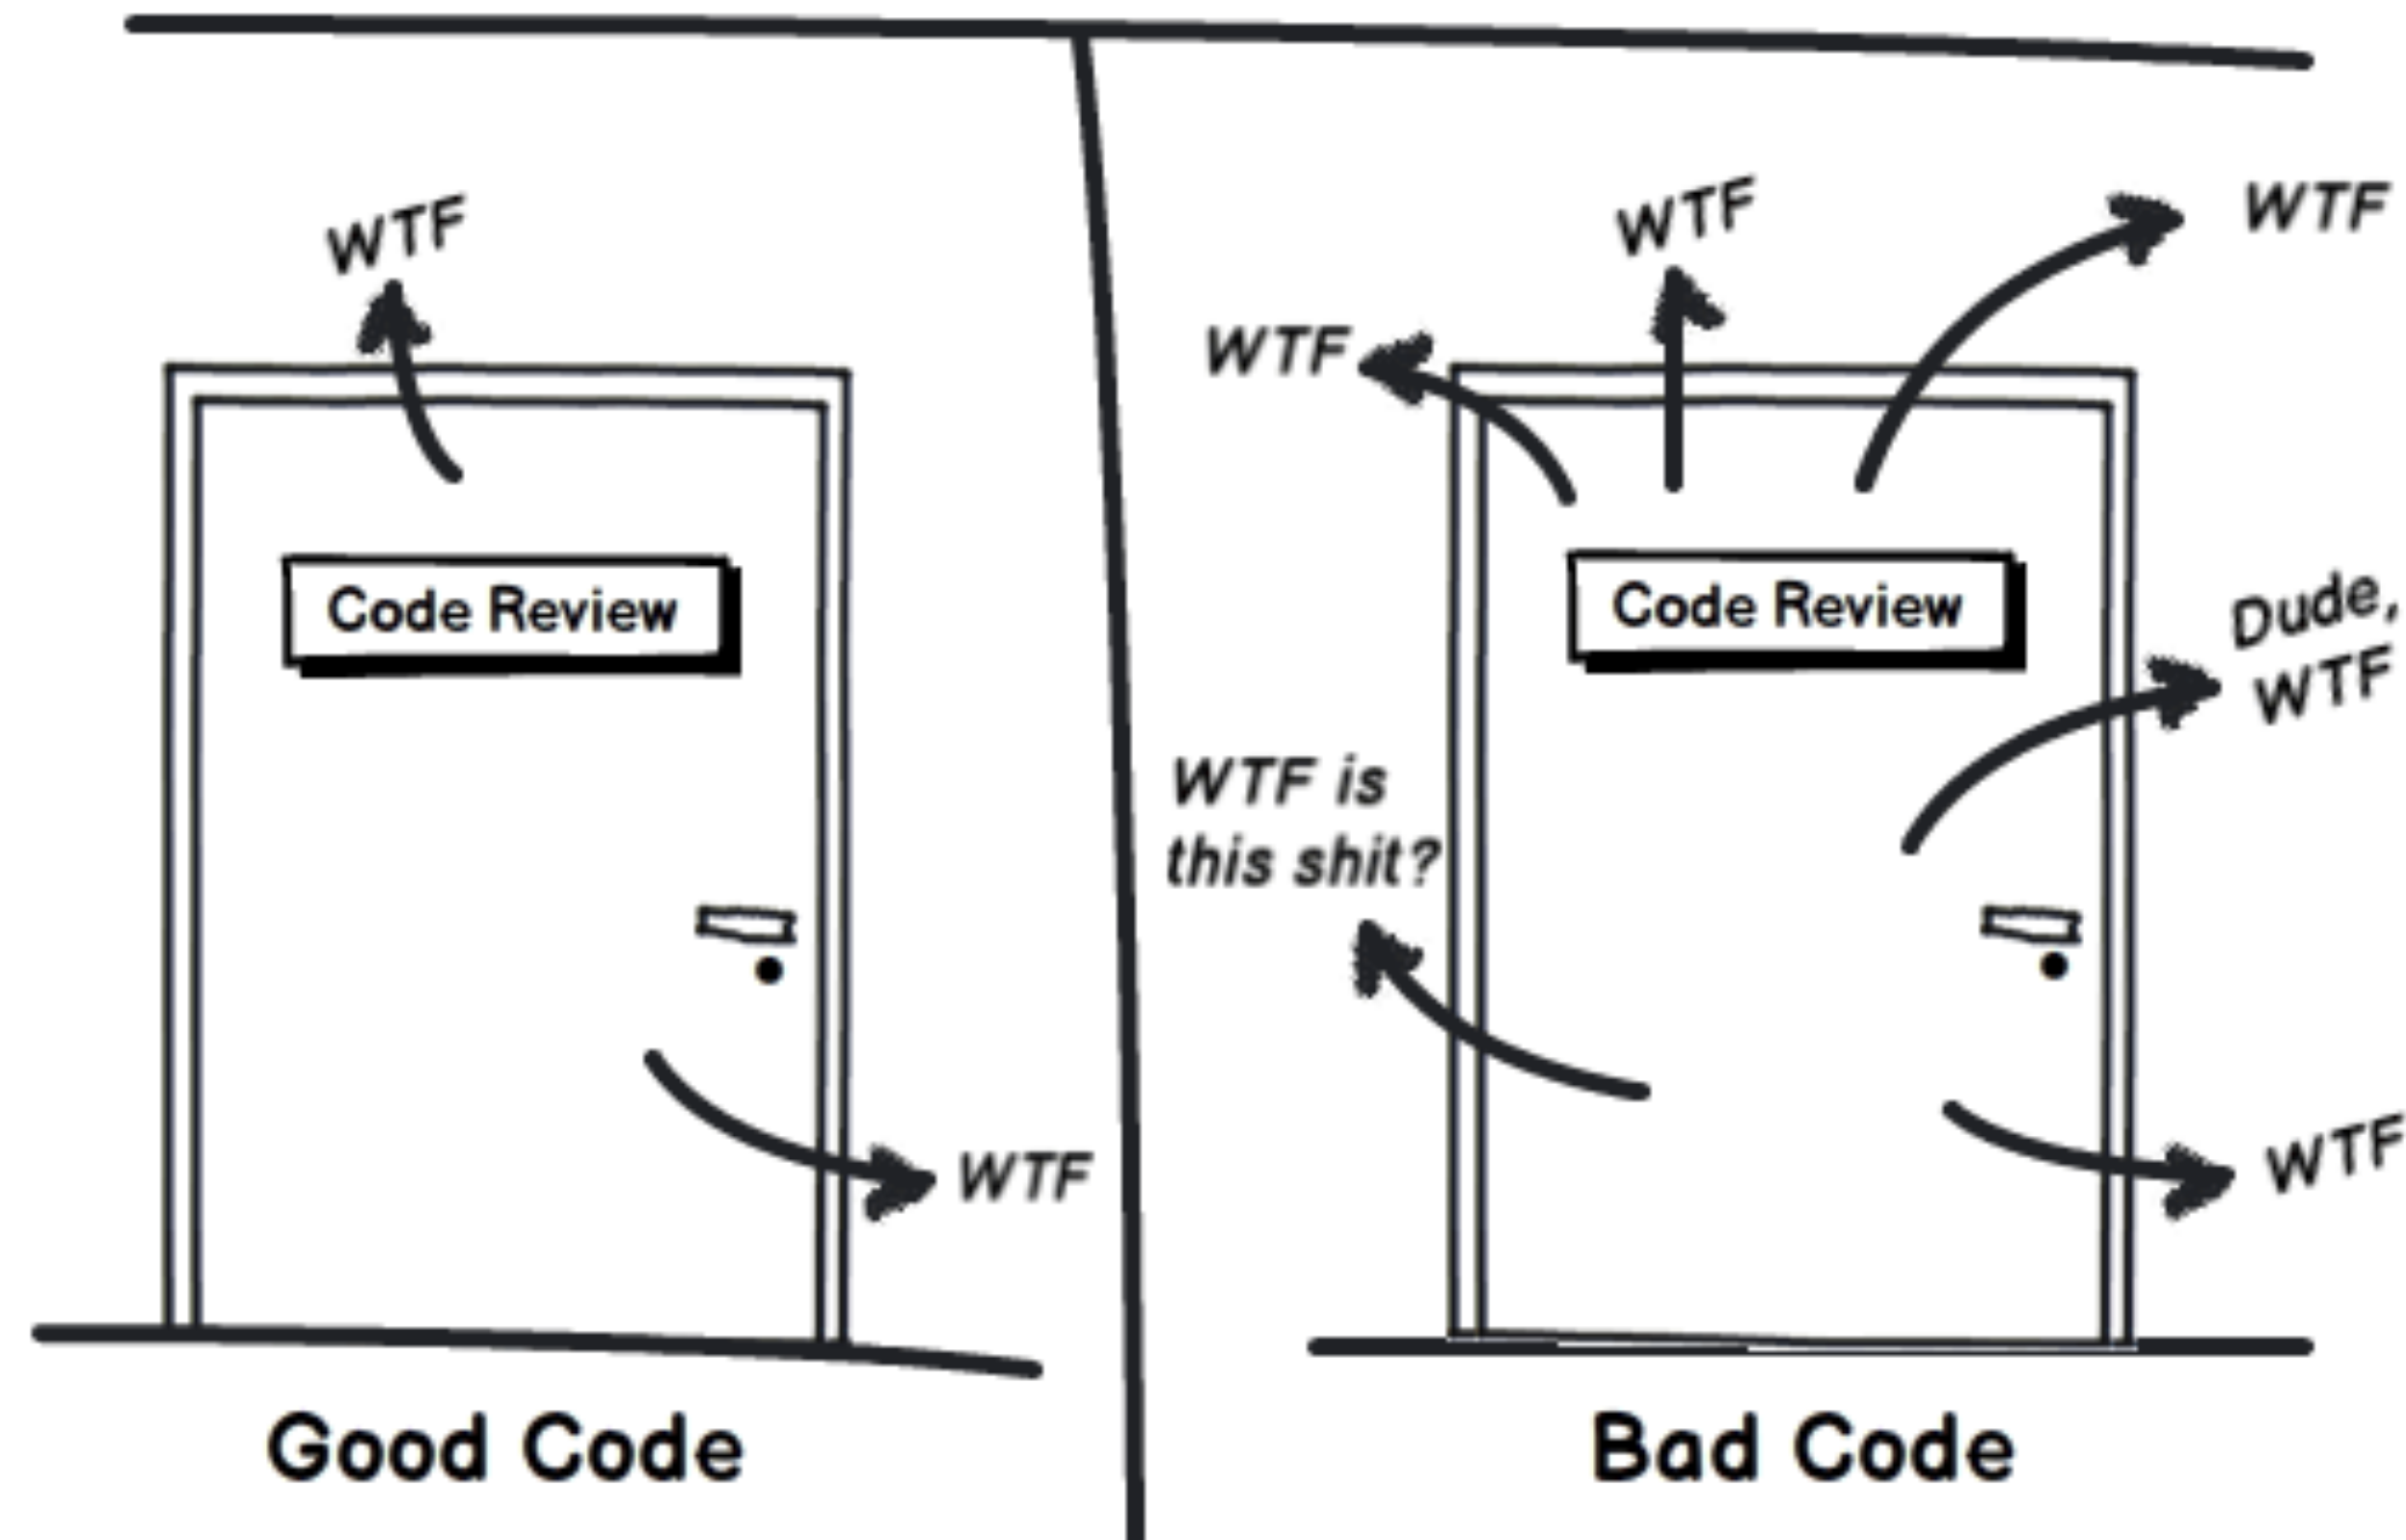
- ▶ **Organisation** Sprawl
- ▶ **Technical** Sprawl
- ▶ Increased **Development Velocity**
- ▶ Building **Trust** and **confidence**

LETS BEGIN ...

#1:

CODE QUALITY

Code Quality Measurement: WTFs/Minute



LINTING/FORMATTING

- ▶ Statically analyse code for lint and format errors
- ▶ Helps write **Idiomatic Go** code
- ▶ Helps Improve readability of the code
- ▶ Makes code **easier to change**
- ▶ Tools used: *"go fmt", "golint", "go vet"*

CODE SMELLS

- ▶ Statically analyse code for *Cyclomatic Complexity*, *DeadCode*, *Duplicated Code*, *ErrCheck* etc
- ▶ Keeps **code quality** in check (**Maintainability**)
- ▶ Helps maintain **Sanity** of the codebase (also of the people reading it)
- ▶ Tools: "*gometalinter*"
- ▶ **Meta Linter** (<https://github.com/alecthomas/gometalinter>)

```
hystrix/hystrix_client_test.go:513:12:warning: ineffectual assignment to err (ineffassign)
hystrix/hystrix_client.go:168::warning: declaration of "err" shadows declaration at hystrix/hystrix_client.go:
163 (vetshadow)
hystrix/hystrix_client.go:178:23:warning: error return value not checked (response.Body.Close()) (errcheck)
hystrix/hystrix_client.go:178::warning: Errors unhandled.,LOW,HIGH (gosec)
hystrix/hystrix_client.go:186::warning: Errors unhandled.,LOW,HIGH (gosec)
hystrix/hystrix_client_test.go:513:12:warning: this value of err is never used (SA4006) (megacheck)
→ heimdall git:(master) x █
```

SECURE CODING

- ▶ Inspect source code for **Security** problems
- ▶ Find vulnerabilities like **SQL Injection, Hardcoded credentials** etc
- ▶ Help write secure code
- ▶ Tools: "**gosec**" (<https://github.com/securego/gosec>)
- ▶ Detects problems with various confidence levels

Available rules

- G101: Look for hard coded credentials
- G102: Bind to all interfaces
- G103: Audit the use of unsafe block
- G104: Audit errors not checked
- G105: Audit the use of math/big.Int.Exp
- G106: Audit the use of ssh.InsecureIgnoreHostKey
- G107: Url provided to HTTP request as taint input
- G201: SQL query construction using format string
- G202: SQL query construction using string concatenation
- G203: Use of unescaped data in HTML templates
- G204: Audit use of command execution
- G301: Poor file permissions used when creating a directory
- G302: Poor file permissions used with chmod
- G303: Creating tempfile using a predictable path
- G304: File path provided as taint input
- G305: File traversal when extracting zip archive
- G401: Detect the usage of DES, RC4, MD5 or SHA1
- G402: Look for bad TLS connection settings
- G403: Ensure minimum RSA key length of 2048 bits
- G404: Insecure random number source (rand)
- G501: Import blacklist: crypto/md5
- G502: Import blacklist: crypto/des
- G503: Import blacklist: crypto/rc4
- G504: Import blacklist: net/http/cgi
- G505: Import blacklist: crypto/sha1

```
[/Users/admin/work/go/src/github.com/gojektech/heimdall/httpclient/client.go:137] - G104: Errors unhandled. (Confidence: HIGH, Severity: LOW)  
> response.Body.Close()
```

```
[/Users/admin/work/go/src/github.com/gojektech/heimdall/httpclient/client.go:145] - G104: Errors unhandled. (Confidence: HIGH, Severity: LOW)  
> _, _ = bodyReader.Seek(0, 0)
```

```
[/Users/admin/work/go/src/github.com/gojektech/heimdall/hystrix/hystrix_client.go:178] - G104: Errors unhandled. (Confidence: HIGH, Severity: LOW)  
> response.Body.Close()
```

```
[/Users/admin/work/go/src/github.com/gojektech/heimdall/hystrix/hystrix_client.go:186] - G104: Errors unhandled. (Confidence: HIGH, Severity: LOW)  
> _, _ = bodyReader.Seek(0, 0)
```

Summary:

Files: 8

Lines: 787

Nosec: 0

Issues: 4

CLOSING COMMENTS ...

- ▶ Make it part of your development process (Makefile)
- ▶ Sample Project (<https://github.com/gojektech/heimdall>)
- ▶ Make it part of your **build pipeline** (**CI/CD**)
- ▶ Changes to the codebase becomes easier
- ▶ Helps maintain **Stability** and **Reliability** of a microservice

#2:

TESTING

**Wrote A Bunch Of
Code Without Testing**



Works


MemeBucket.com

UNIT TESTING

- ▶ Unit tests test a unit/function (Individual component)
- ▶ These are run very often during development
- ▶ Are very large in number (# of unit tests)
- ▶ Other practices include **TDD (Test Driven Development)**
- ▶ Tools: "in-built **Go testing framework**"
- ▶ Others (**Testify**): (<https://github.com/stretchr/testify>)

UNIT TESTS SUFFICIENT ?

passing unit tests





Guillaume Malette @gmalette

1h

Replying to @iamdeveloper

Running integration tests:

- Ship still afloat after four watertight compartments flooded: PASS
- Band still plays after collision with iceberg: PASS

1 1 1

INTEGRATION TESTING

- ▶ Integration tests help in testing multiple components working together
- ▶ Cover **End to end** flows
- ▶ Lesser in number than unit tests
- ▶ Helps in finding regressions in the codebase effectively
- ▶ Tools: "**godog**" (BDD)
- ▶ Link: (<https://github.com/DATA-DOG/godog>)

Feature: customer legacy login

In order to access application resources

As a Customer

I need to login with my credentials

Scenario: Successful login


Given correct credentials for login

When customer logs in

Then the login response code should be 201

And the login response should have valid details

TEST PYRAMID



LOAD AND PERFORMANCE TESTING

- ▶ Test behaviour of app under various loads
- ▶ Helps find performance bottlenecks
- ▶ Also Memory Leaks
- ▶ Helps with **Capacity Planning** and finding **SLO**
- ▶ Sizing: **Step Load, Linear Load, Spiky load, Soak test**
- ▶ Tools: **"Gatling"**
- ▶ In go: **"Vegeta"** (<https://github.com/tsenart/vegeta>)



50 rps, 1 min, latency on both services

Requests	[total, rate]	3000, 50.01
Duration	[total, attack, wait]	1m0.217022968s, 59.983753s, 233.269968ms
Latencies	[mean, 50, 95, 99, max]	190.12369ms, 186.887664ms, 257.013489ms, 269.958324ms, 348.796101ms
Bytes In	[total, mean]	527506, 175.84
Bytes Out	[total, mean]	348000, 116.00
Success	[ratio]	99.57%
Status Codes	[code:count]	201:2987 500:13

OTHERS

- ▶ Contract Testing
- ▶ Chaos Testing

CLOSING COMMENTS ...

- ▶ Integral part of **Development** process
- ▶ Integral part of your **CI/CD pipeline**
- ▶ Helps build **confidence** in microservice and the codebase
- ▶ Helps move faster (Increases Developer velocity)
- ▶ Helps with **Scalability, Performance, Stability and Reliability**

#3:

OBSERVABILITY



Charity Majors

@mipsytipsy

Follow



Replying to @mipsytipsy @mattklein123

Observability, otoh, is about being able to understand the inner workings of your software and systems by asking questions and observing the answers on the outside. Any question — no particular bias toward actionable alerts or problems.

OBSERVABILITY BROADLY INCLUDES:

MONITORING

LOGGING

~~DISTRIBUTED TRACING~~

MONITORING

**Monitoring is the act of checking
the behaviour and outputs of a
system and its components**

MONITORING

- ▶ It is all about looking out for **presence/absence of patterns**
- ▶ Is used to report **overall health** of the system
- ▶ Includes **Key Business, Microservice and System Level Metrics**
- ▶ Think **USE** (Utilisation, Saturation and Errors)
- ▶ Think **RED** (Response times, errors, duration)

KEY METRICS

- ▶ Host and Infrastructure metrics
 - ▶ CPU
 - ▶ RAM
 - ▶ Threads
 - ▶ File Descriptors
 - ▶ Database connections

KEY METRICS

- ▶ Microservice Key Metrics
 - ▶ Availability
 - ▶ SLA
 - ▶ Latency
 - ▶ Success
 - ▶ Errors


KEY METRICS

- ▶ Business Metrics
 - ▶ # of bookings completed
 - ▶ # of successful logins
 - ▶ # of available drivers
 - ▶ # of booking cancelations

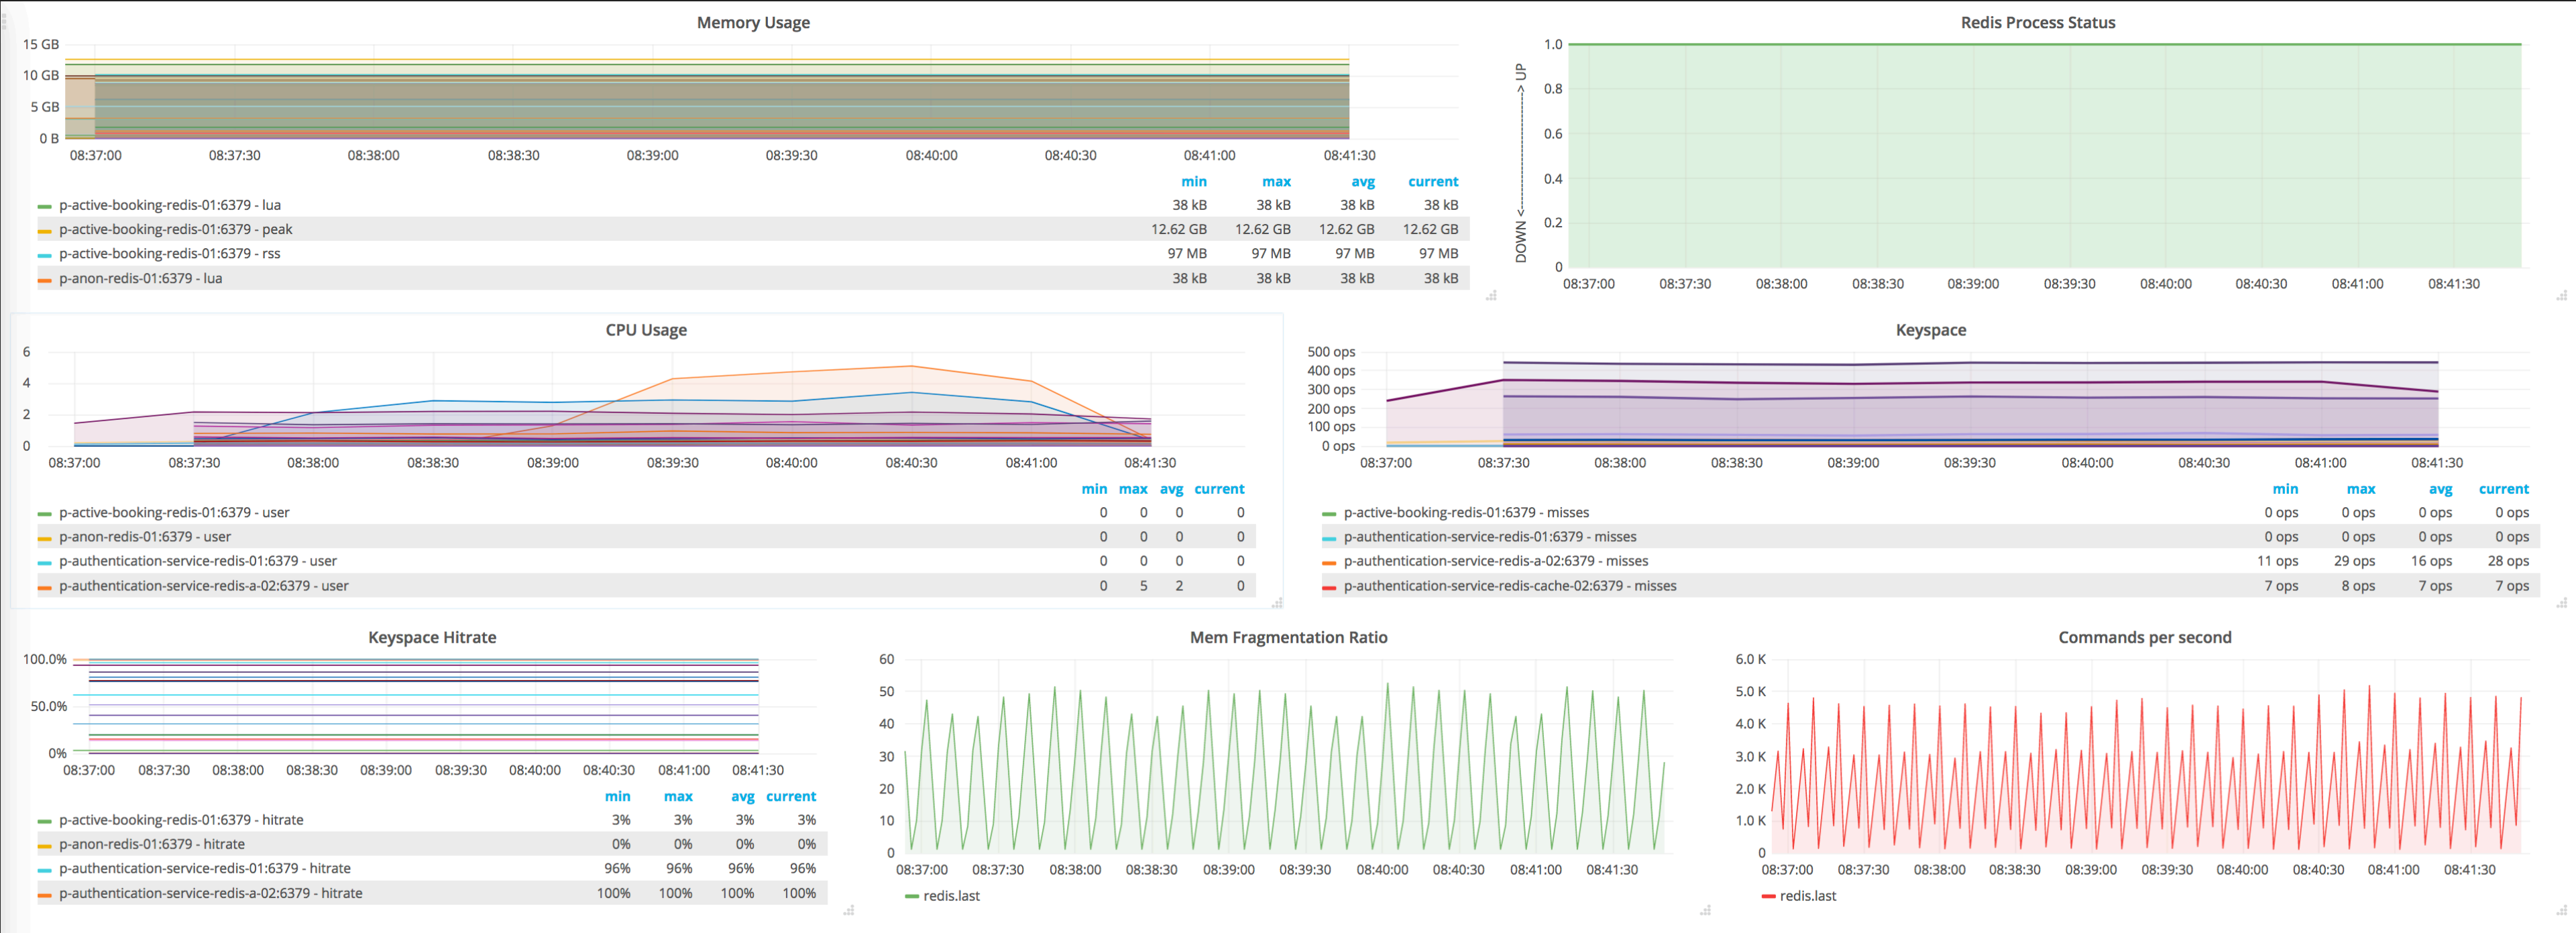
DASHBOARD

- ▶ Reflect the state and **health of your system**
- ▶ Should capture key metrics
- ▶ **BusinessMetrics, UptimeMetrics, SystemMetrics**
- ▶ Grafana/NewRelic at GO-JEK

APM (APPLICATION PERFORMANCE MONITORING)

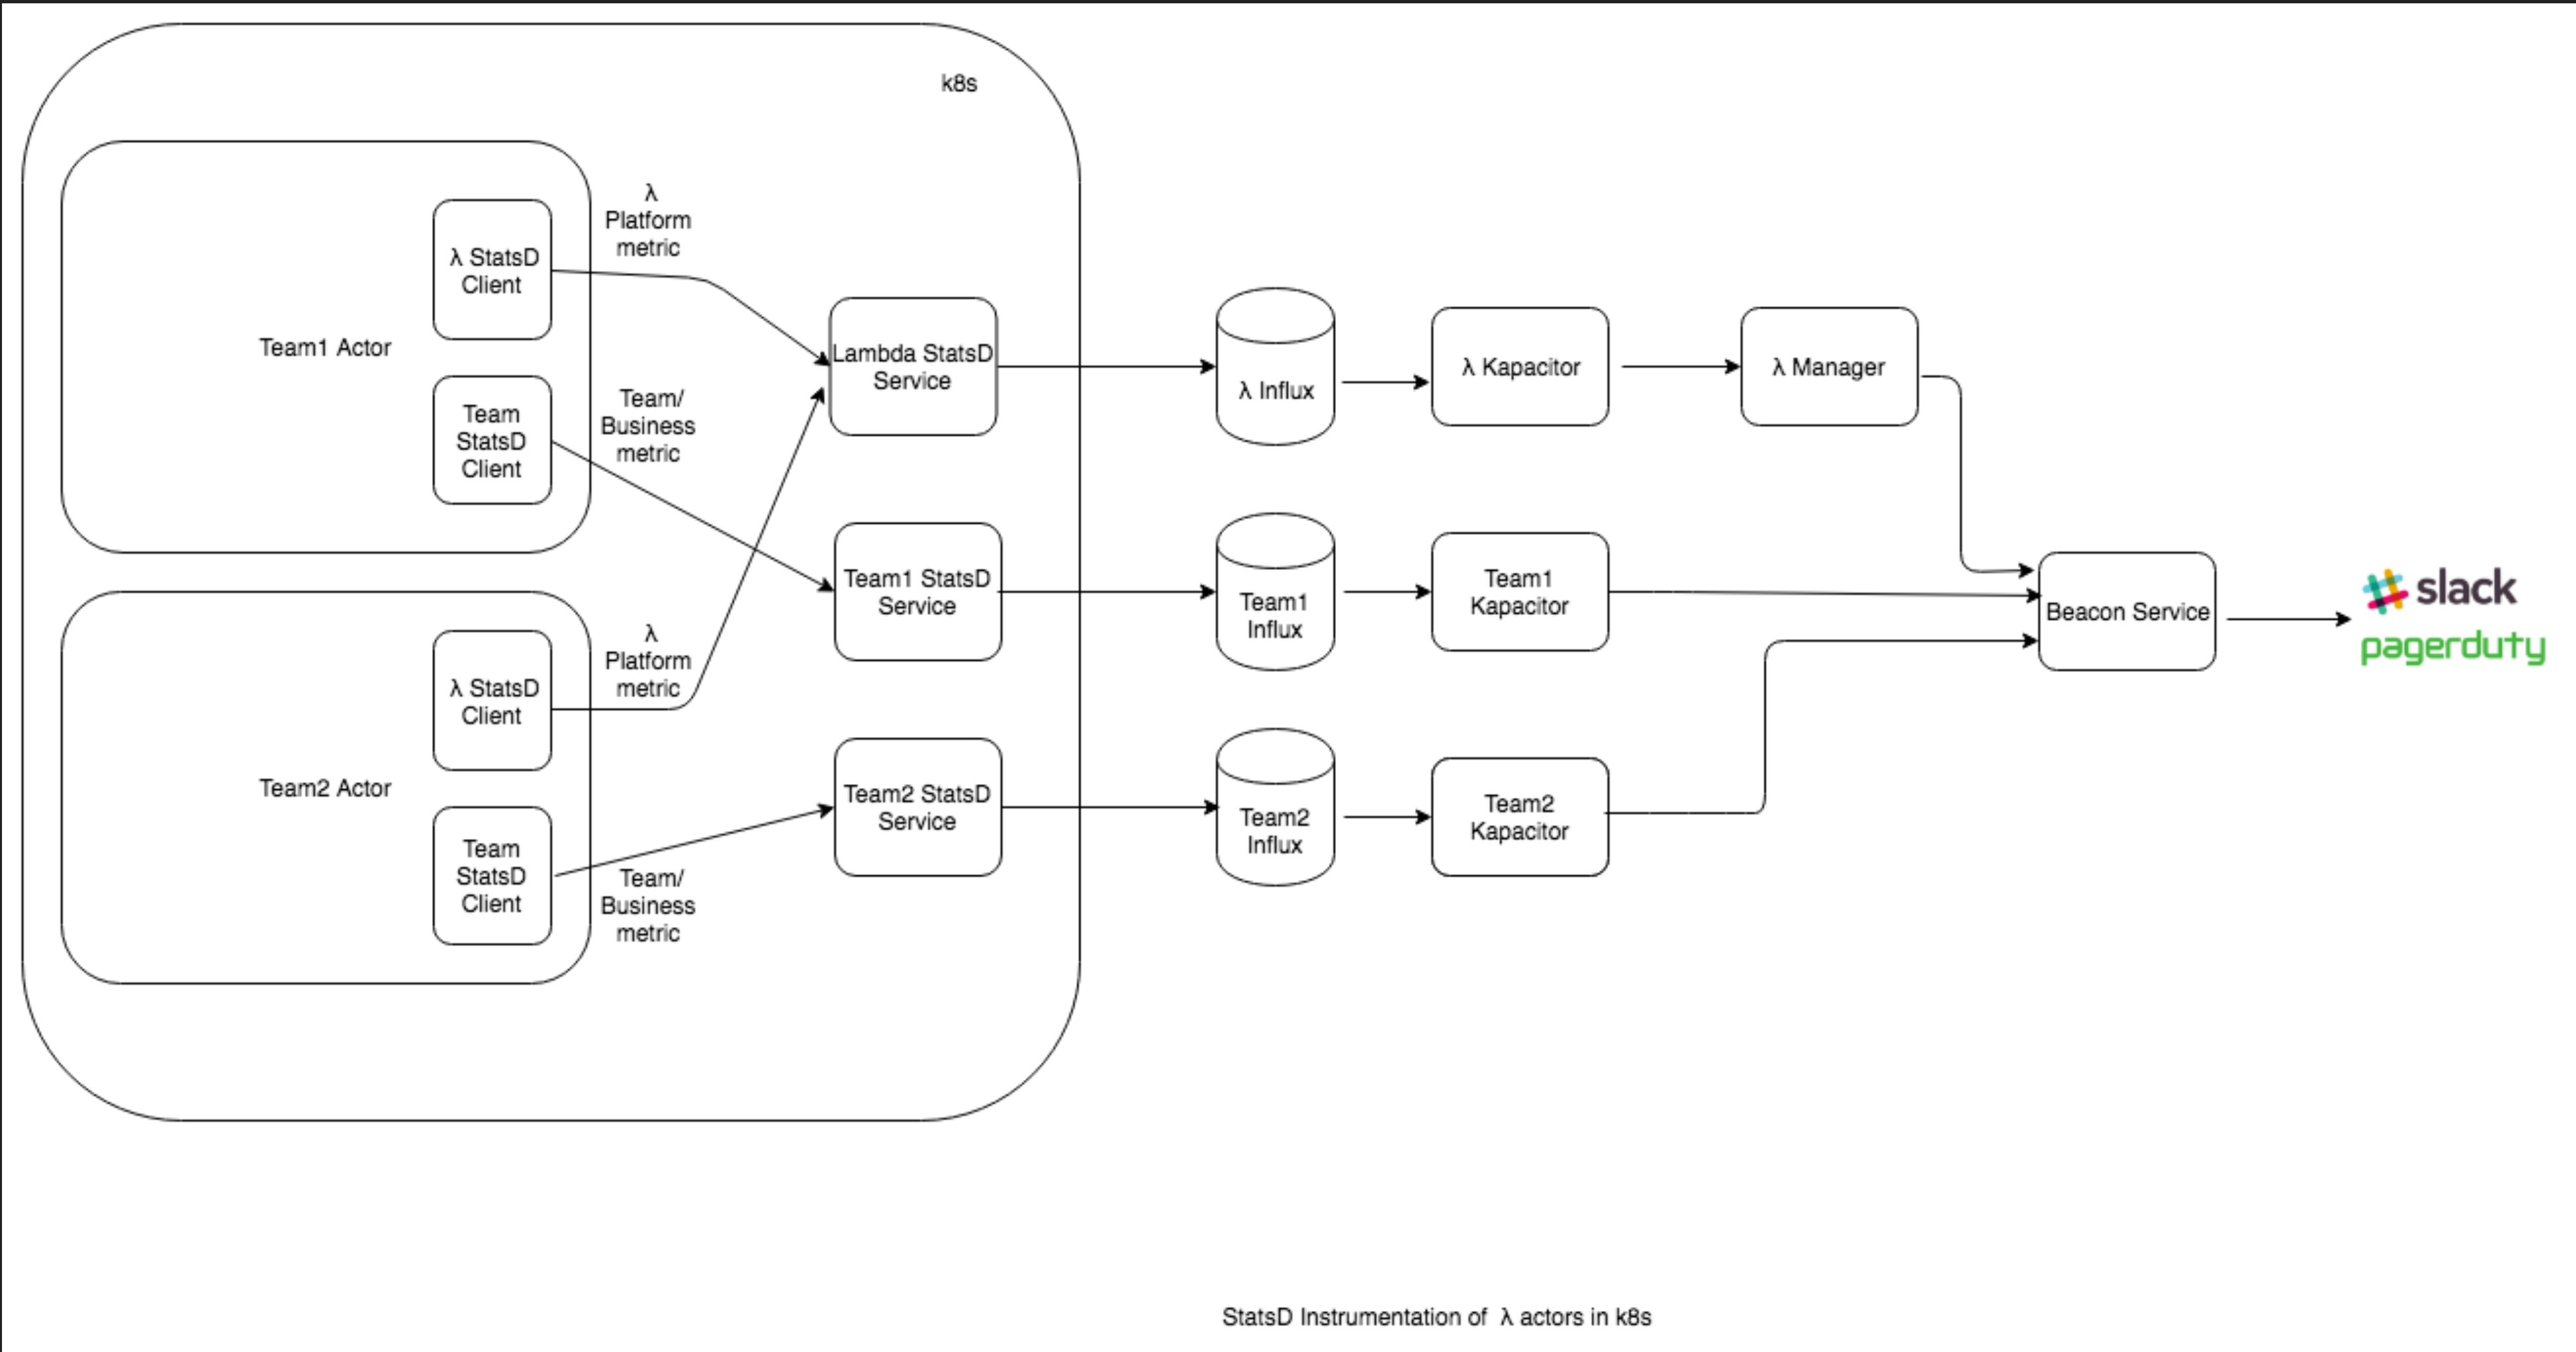


SYSTEM LEVEL METRICS (GRAFANA) – TICK STACK




ALERTING

- ▶ Notifying team when an anomaly in key metrics is detected
- ▶ **Thresholds for various key metrics** is set for alerting
- ▶ Alerts should be **actionable**
- ▶ Help in alerting teams and recovering before a catastrophe
- ▶ **TICK** stack (Telegraf, Influx, Chronograph, Kapacitor)



LOGGING

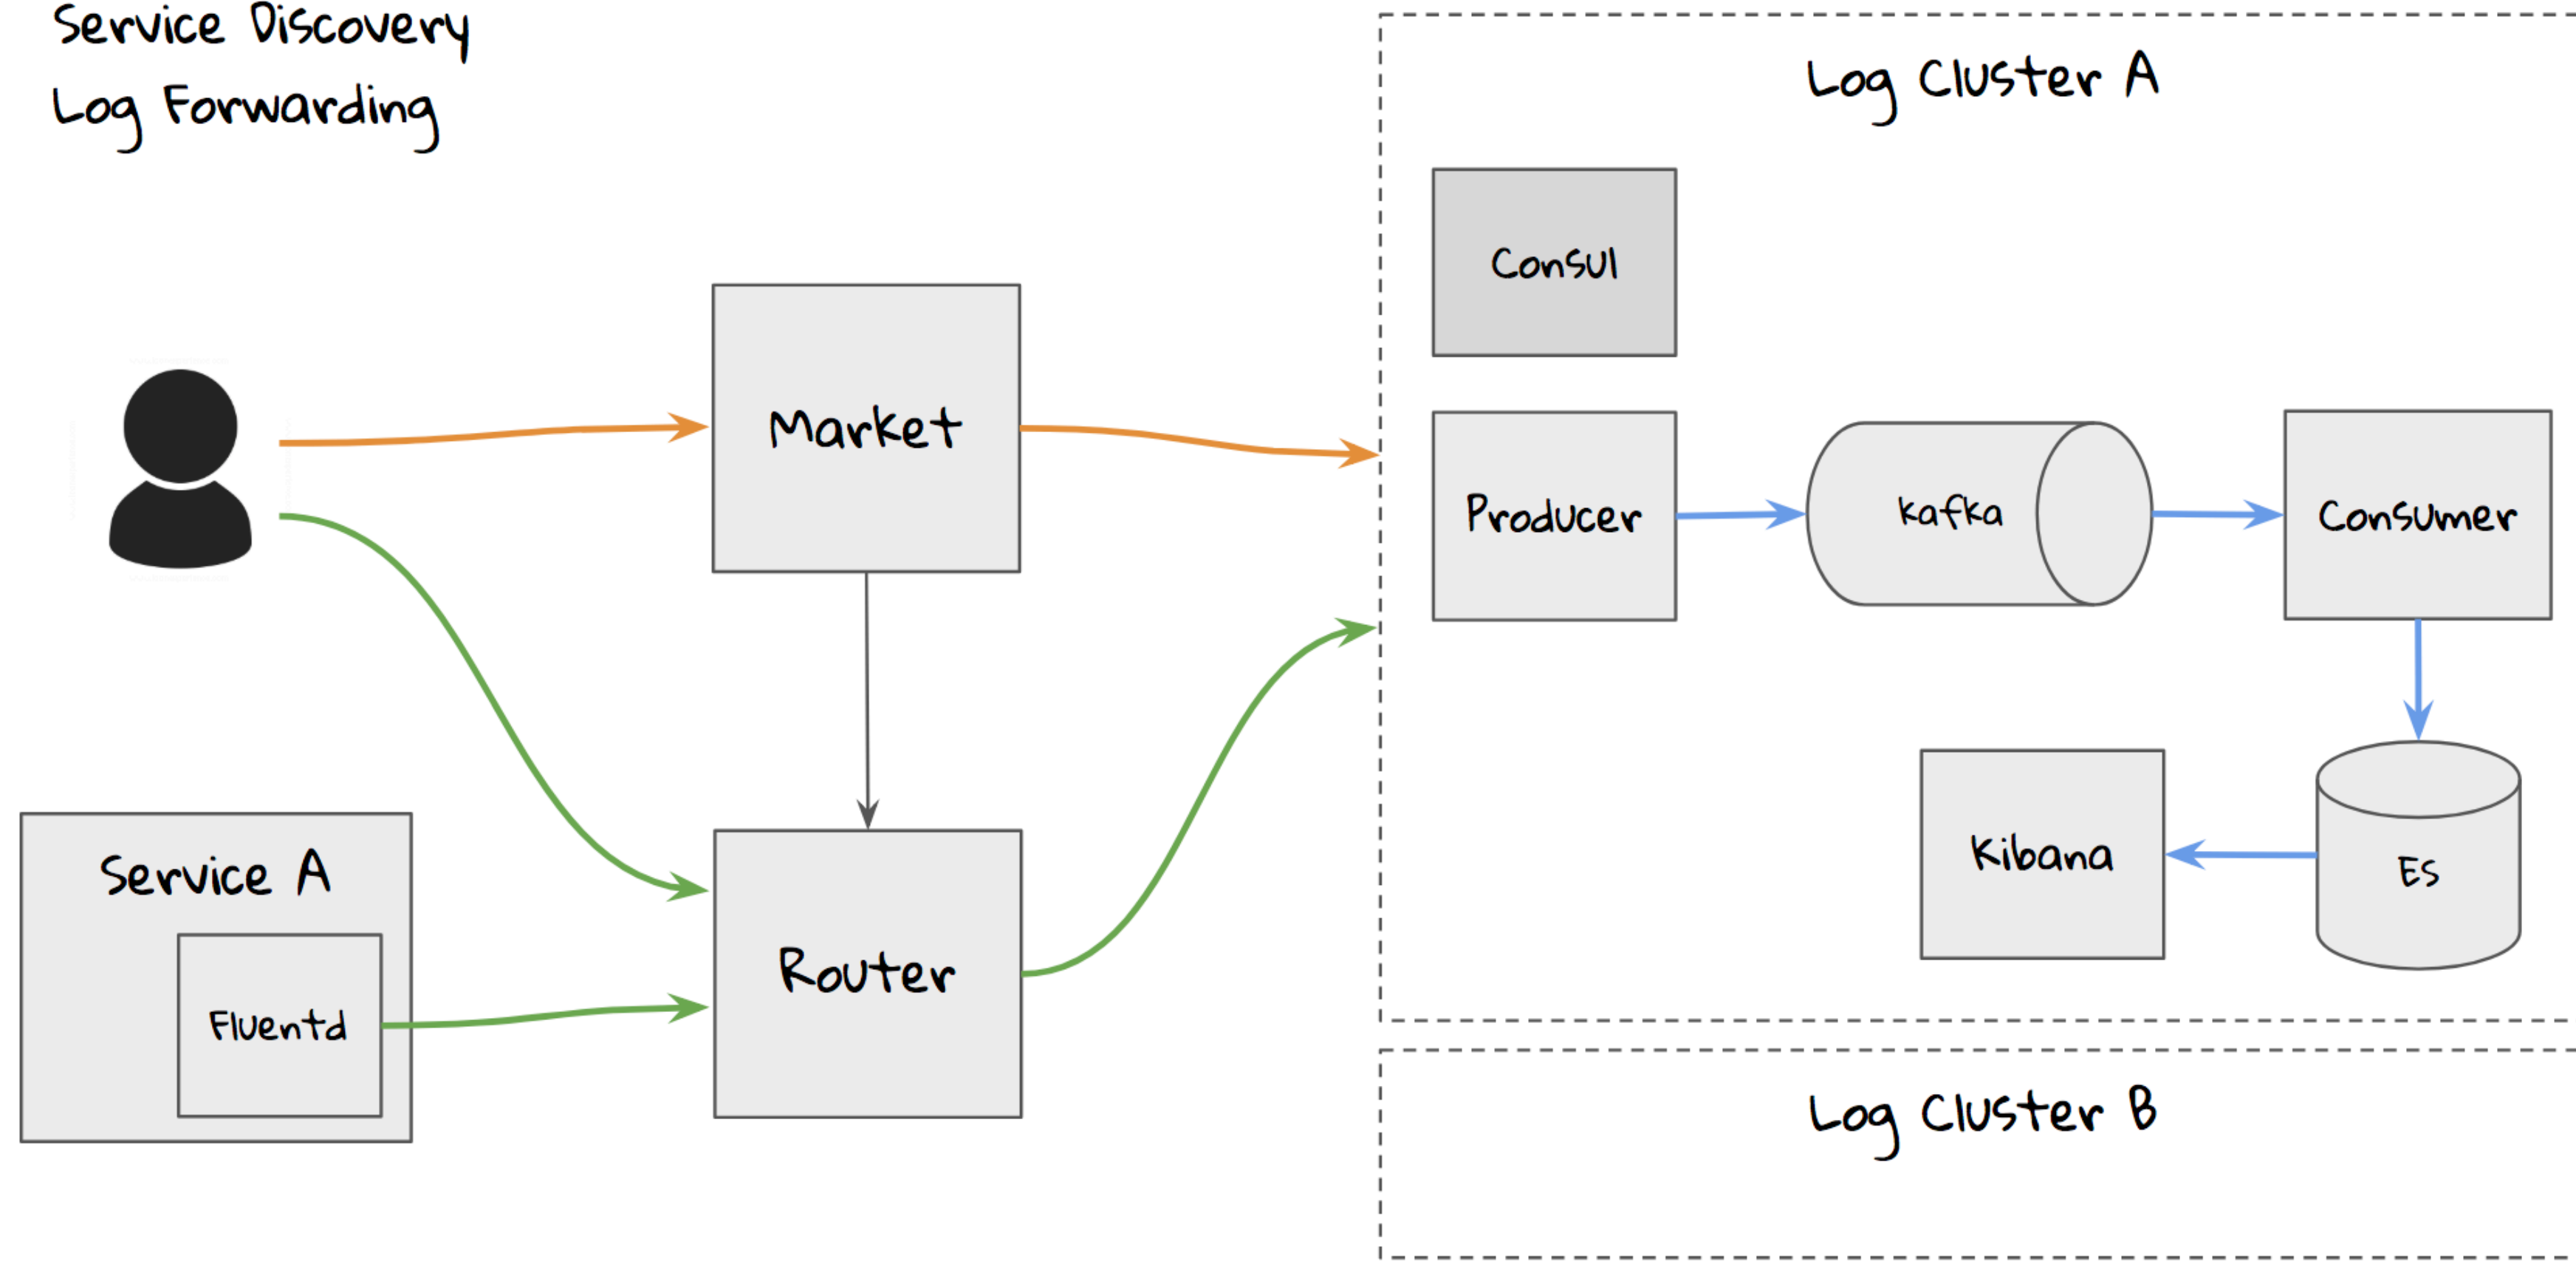


LOGGING

- ▶ Logging helps ***describe system state*** at any point in time
- ▶ Helps in debugging problems with the system
- ▶ Structured logging
- ▶ **Log centralisation, aggregation and analysis** is the key
- ▶ **Barito-Log** (github.com/BaritoLog)

Barito Overview

- Provisioning
- Service Discovery
- Log Forwarding



INSTRUMENTATION OF CODE

- ▶ Ability to measure various aspects of system behaviour
- ▶ Code needs to be written to monitor specific components
- ▶ It includes **Structured logging**, **StatsD metrics** (Counters, Gauges, Histograms) and **Error/Exception tracking (Sentry)**
- ▶ Should also be able to capture **Application key performance metrics (NewRelic)**

```
type StatsD struct {
    client *statsd.Client
}

func (reporter *StatsD) Incr(key string, tags []string) {
    if reporter.client != nil {
        reporter.client.Incr(key, tags, 1)
    }
}

reporter.Incr("customers.login.count", []string{"login"})
```

CLOSING COMMENTS ...

- ▶ Insights into **Service Health** and **Behaviour**
- ▶ Insights into **Service Performance**
- ▶ Helps providing **Debuggability in Production**
- ▶ Help bring visibility into system
- ▶ Knowing when things go wrong
- ▶ Helps with **Stability** of the system

OTHERS

- ▶ Resiliency Patterns ([Gophercon India 2018 Talk](#))
- ▶ Security
- ▶ Documentation
- ▶ Canarying (Testing in production)
- ▶ CI/CD (Stable deployment process)

IN THE END ...

WRAPPING UP

- ▶ **Code quality**
- ▶ **Testing**
- ▶ **Observability**

TAKE AWAYS

- ▶ Code Health Metrics dashboard
- ▶ **Observability** driven Development
- ▶ Tests as **safety harness** - Remember Test Pyramid

IN CONCLUSION...

STANDARDISING

QUALITY

IS THE GOAL

AVAILABILITY
IS THE GOAL

WHAT NEXT ?

- ▶ **Automating** the process of production readiness
- ▶ Defining a **production readiness score**
- ▶ Assigning every microservice a **readiness score**
- ▶ Help teams improve quality of microservices in Production
- ▶ **Measure** the impact and **improve**

REFERENCES

- [Production Ready Microservices - Susan Fowler](#)
- [Google SRE Book](#)
- [Microservices Standardisation](#)
- [Resiliency in Distributed Systems](#)
- [TICK stack](#)
- [SLA](#)

THANKS FOR
LISTENING

QUESTIONS?